

End-Semester Examination

Allotted time: 60 minutes

Marks: 32

Instructions:

- There are four questions with sub-parts, printed over two sides of a sheet. Please check if your sheet has all the questions. Marks will be normalized by a factor of 10/32.
- Discussions amongst the students are not allowed. Any dishonesty shall be penalized heavily.
- Be clear in your arguments. Vague arguments shall not be given full credit.
- Please note that in class we discussed NASM x86 assembly and in assignments we discussed GNU x86 assembly. The main difference is in the order of operands. NASM uses destination first but GNU Assembly does not use destination as a first operand.

1. Consider a 5-bit floating-point representation based on the IEEE floating-point format, with one sign bit, two exponent bits ($k = 2$), and two fraction bits ($n = 2$).

(a) Compute the bias.

(b) What is the binary representation of $-\frac{31}{8}$ and $\frac{56}{32}$ in the above floating point representation.

[2+3+3 marks]

2. Explain the given GNU x86 assembly snippet.

```
push %ebp  
mov %esp,%ebp  
mov 0x8(%ebp),%eax  
# lea (%esi,%ebx,4), %edx corresponds to setting %edx to the address %esi+%ebx*4  
lea (%eax,%eax,4),%eax  
lea 0xffffffff(%eax,%eax,4),%eax  
mov %ebp,%esp  
pop %ebp  
ret
```

[8 marks]

3. The 5 stages of the processor have the following latencies:

	Fetch	Decode	Execute	Memory	Writeback
a.	300ps	400ps	350ps	550ps	100ps
b.	200ps	150ps	100ps	190ps	140ps

Assume that when pipelining, each pipeline stage costs 20ps extra for the registers between pipeline stages.

- (a) Non-pipelined processor: what is the cycle time? What is the latency of an instruction? What is the throughput?
- (b) Fully Pipelined processor: What is the cycle time? What is the latency of an instruction? What is the throughput?
- (c) If you could split one of the pipeline stages into 2 equal halves, which one would you choose? What is the new cycle time? What is the new latency? What is the new throughput?
4. (a) For a direct-mapped cache design with a 32-bit address and byte-addressable memory, the following bits of the address are used to access the cache:

	Tag	Index	Offset
a.	31-10	9-5	4-0
b.	31-12	11-6	5-0

For each configuration (a and b):

- i. What is the cache block size (in words)?
- ii. How many entries (blocks) does the cache have?

- (b) You have a 2-way set associative L1 cache that is 8KB, with 4-word cache lines. Writing and reading data to L2 takes 10 cycles. You get the following sequence of writes to the cache – each is a 32-bit address in hexadecimal:

0x1000

0x1004

0x1010

0x11c0

0x2000

0x21c0

0x3400

0x3404

0x3f00

0x2004

0x1004

- i. How many cache misses occur with an LRU policy?
- ii. How many cache misses occur with a FIFO policy?
- iii. Would the miss-rate increase or decrease if the cache was the same size, but direct-mapped (in above case)? Explain.

[2+3+3 marks]

[1.5+1.5 marks]

[2+1+2 marks]

x86-64 Reference Sheet (GNU assembler format)

Instructions

Data movement

`movq Src, Dest` Dest = Src
`movsbq Src, Dest` Dest (quad) = Src (byte), sign-extend
`movzbq Src, Dest` Dest (quad) = Src (byte), zero-extend

Conditional move

`cmove Src, Dest` Equal / zero
`cmovev Src, Dest` Not equal / not zero
`cmove Src, Dest` Negative
`cmovea Src, Dest` Nonnegative
`cmoveg Src, Dest` Greater (signed >)
`cmovege Src, Dest` Greater or equal (signed ≥)
`cmovel Src, Dest` Less (signed <)
`cmovele Src, Dest` Less or equal (signed ≤)
`cmovea Src, Dest` Above (unsigned >)
`cmoveae Src, Dest` Above or equal (unsigned ≥)
`cmoveb Src, Dest` Below (unsigned <)
`cmovebe Src, Dest` Below or equal (unsigned ≤)

Control transfer

`cmpq Src2, Src1` Sets CCs Src1 Src2
`testq Src2, Src1` Sets CCs Src1 & Src2
`jmp label` jump
`je label` jump equal
`jne label` jump not equal
`js label` jump negative
`jns label` jump non-negative
`jg label` jump greater (signed >)
`jge label` jump greater or equal (signed ≥)
`jl label` jump less (signed <)
`jle label` jump less or equal (signed ≤)
`ja label` jump above (unsigned >)
`jb label` jump below (unsigned <)
`pushq Src` $\%rsp = \%rsp - 8$, $Mem[\%rsp] = Src$
`popq Dest` Dest = $Mem[\%rsp]$, $\%rsp = \%rsp + 8$
`call label` push address of next instruction, `jmp label`
`ret` $\%rip = Mem[\%rsp]$, $\%rsp = \%rsp + 8$

Arithmetic operations

`leaq Src, Dest` Dest = address of Src
`incq Dest` Dest = Dest + 1
`decq Dest` Dest = Dest - 1
`addq Src, Dest` Dest = Dest + Src
`subq Src, Dest` Dest = Dest - Src
`imulq Src, Dest` Dest = Dest * Src
`xorq Src, Dest` Dest = Dest ^ Src
`orq Src, Dest` Dest = Dest | Src
`andq Src, Dest` Dest = Dest & Src
`negq Dest` Dest = - Dest
`notq Dest` Dest = ~ Dest
`salq k, Dest` Dest = Dest ≪ k
`sarq k, Dest` Dest = Dest ≫ k (arithmetic)
`shrq k, Dest` Dest = Dest ≫ k (logical)

Addressing modes

- Immediate
 $\$val Val$
 val : constant integer value
`movq $7, %rax`
- Normal
 $(R) Mem[Reg[R]]$
 R : register R specifies memory address
`movq (%rcx), %rax`
- Displacement
 $D(R) Mem[Reg[R]+D]$
 R : register specifies start of memory region
 D : constant displacement D specifies offset
`movq 8(%rdi), %rdx`
- Indexed
 $D(Rb,Ri,S) Mem[Reg[Rb]+S*Reg[Ri]+D]$
 D : constant displacement 1, 2, or 4 bytes
 Rb : base register: any of 8 integer registers
 Ri : index register: any, except %esp
 S : scale: 1, 2, 4, or 8
`movq 0x100(%rcx,%rax,4), %rdx`

Instruction suffixes

b byte
w word (2 bytes)
l long (4 bytes)
q quad (8 bytes)

Condition codes

CF Carry Flag
ZF Zero Flag
SF Sign Flag
OF Overflow Flag

Integer registers

<code>%rax</code>	Return value
<code>%rbx</code>	Callee saved
<code>%rcx</code>	4th argument
<code>%rdx</code>	3rd argument
<code>%rsi</code>	2nd argument
<code>%rdi</code>	1st argument
<code>%rbp</code>	Callee saved
<code>%rsp</code>	Stack pointer
<code>%r8</code>	5th argument
<code>%r9</code>	6th argument
<code>%r10</code>	Scratch register
<code>%r11</code>	Scratch register
<code>%r12</code>	Callee saved
<code>%r13</code>	Callee saved
<code>%r14</code>	Callee saved
<code>%r15</code>	Callee saved

CS107 x86-64 Reference Sheet

Common instructions

mov src, dst	dst = src
movsbl src, dst	byte to int, sign-extend
movzbl src, dst	byte to int, zero-fill
cmove src, reg	reg = src when condition holds, using same condition suffixes as jmp
lea addr, dst	dst = addr
add src, dst	dst += src
sub src, dst	dst -= src
imul src, dst	dst *= src
neg dst	dst = -dst (arith inverse)
imulq S	signed full multiply R[%rdx]:R[%rax] <- S * R[%rax]
mulq S	unsigned full multiply same effect as imulq
idivq S	signed divide R[%rdx] <- R[%rdx]:R[%rax] mod S R[%rax] <- R[%rdx]:R[%rax] / S
divq S	unsigned divide - same effect as idivq
cqto	R[%rdx]:R[%rax] <- SignExtend(R[%rax])
sal count, dst	dst <= count
sar count, dst	dst >= count (arith shift)
shr count, dst	dst >= count (logical shift)
and src, dst	dst &= src
or src, dst	dst = src
xor src, dst	dst ^= src
not dst	dst = ~dst (bitwise inverse)
cmp a, b	b-a, set flags
test a, b	a&b, set flags
set dst	sets byte at dst to 1 when condition holds, 0 otherwise, using same condition suffixes as jmp
jmp label	jump to label (unconditional)
je label	jump equal ZF=1
jne label	jump not equal ZF=0
js label	jump negative SF=1
jns label	jump not negative SF=0
jg label	jump > (signed) ZF=0 and SF=OF
jge label	jump >= (signed) SF=OF
jl label	jump < (signed) SF!=OF
jle label	jump <= (signed) ZF=1 or SF!=OF
ja label	jump > (unsigned) CF=0 and ZF=0
jae label	jump >= (unsigned) CF=0
jb label	jump < (unsigned) CF=1
jbe label	jump <= (unsigned) CF=1 or ZF=1

push src	add to top of stack Mem[-%rsp] = src
pop dst	remove top from stack dst = Mem[%rsp++]
call fn	push %rip, jmp to fn
ret	pop %rip

Condition codes/flags

ZF	Zero flag
SF	Sign flag
CF	Carry flag
OF	Overflow flag

Addressing modes

Example source operands to **mov**

Immediate

mov \$0x5, dst

\$val
source is constant value

Register

mov %rax, dst

%R
R is register
source in %R register

Direct

mov 0x4033d0, dst

0xaddr
source read from Mem[0xaddr]

Indirect

mov (%rax), dst

(%R)
R is register
source read from Mem[%R]

Indirect displacement

mov 8(%rax), dst

D(%R)
R is register
D is displacement
source read from Mem[%R + D]

Indirect scaled-index

mov 8(%rsp, %rcx, 4), dst

D(%RB,%RI,S)
RB is register for base
RI is register for index (0 if empty)
D is displacement (0 if empty)
S is scale 1, 2, 4 or 8 (1 if empty)
source read from:
Mem[%RB + D + S*%RI]

CS107 x86-64 Reference Sheet

Registers

%rip	Instruction pointer
%rsp	Stack pointer
%rax	Return value
%rdi	1st argument
%rsi	2nd argument
%rdx	3rd argument
%rcx	4th argument
%r8	5th argument
%r9	6th argument
%r10, %r11	Callee-owned
%rbx, %rbp, %r12-%15	Caller-owned

Instruction suffixes

b	byte
w	word (2 bytes)
l	long /doubleword (4 bytes)
q	quadword (8 bytes)

Suffix is elided when can be inferred from operands. e.g. operand %rax implies q, %eax implies l, and so on

Register Names

64-bit register	32-bit sub-register	16-bit sub-register	8-bit sub-register
%rax	%eax	%ax	%al
%rbx	%ebx	%bx	%bl
%rcx	%ecx	%cx	%cl
%rdx	%edx	%dx	%dl
%rsi	%esi	%si	%sil
%rdi	%edi	%di	%dil
%rbp	%ebp	%bp	%bp1
%rsp	%esp	%sp	%sp1
%r8	%r8d	%r8w	%r8b
%r9	%r9d	%r9w	%r9b
%r10	%r10d	%r10w	%r10b
%r11	%r11d	%r11w	%r11b
%r12	%r12d	%r12w	%r12b
%r13	%r13d	%r13w	%r13b
%r14	%r14d	%r14w	%r14b
%r15	%r15d	%r15w	%r15b

NASM Intel x86 Assembly Language Cheat Sheet

Instruction	Effect	Examples
Copying Data		
<code>mov dest,src</code>	Copy src to dest	<code>mov eax,10</code> <code>mov eax,[2000]</code>
Arithmetic		
<code>add dest,src</code>	<code>dest = dest + src</code>	<code>add esi,10</code>
<code>sub dest,src</code>	<code>dest = dest - src</code>	<code>sub eax, ebx</code>
<code>mul reg</code>	<code>edx:eax = eax * reg</code>	<code>mul esi</code>
<code>div reg</code>	<code>edx = edx:eax mod reg</code> <code>eax = edx:eax ÷ reg</code>	<code>div edi</code>
<code>inc dest</code>	Increment destination	<code>inc eax</code>
<code>dec dest</code>	Decrement destination	<code>dec word [0x1000]</code>
Function Calls		
<code>call label</code>	Push eip, transfer control	<code>call format_disk</code>
<code>ret</code>	Pop eip and return	<code>ret</code>
<code>push item</code>	Push item (constant or register) to stack. I.e.: <code>esp=esp-4; memory[esp] = item</code>	<code>push dword 32</code> <code>push eax</code>
<code>pop [reg]</code>	Pop item from stack and store to register I.e.: <code>reg=memory[esp]; esp=esp+4</code>	<code>pop eax</code>
Bitwise Operations		
<code>and dest, src</code>	<code>dest = src & dest</code>	<code>and ebx, eax</code>
<code>or dest,src</code>	<code>dest = src dest</code>	<code>or eax,[0x2000]</code>
<code>xor dest, src</code>	<code>dest = src ^ dest</code>	<code>xor ebx, 0xffffffff</code>
<code>shl dest,count</code>	<code>dest = dest << count</code>	<code>shl eax, 2</code>
<code>shr dest,count</code>	<code>dest = dest >> count</code>	<code>shr dword [eax],4</code>
Conditionals and Jumps		
<code>cmp b,a</code>	Compare b to a; must immediately precede any of the conditional jump instructions	<code>cmp eax,0</code>
<code>je label</code>	Jump to label if <code>b == a</code>	<code>je endloop</code>
<code>jne label</code>	Jump to label if <code>b != a</code>	<code>jne loopstart</code>
<code>jg label</code>	Jump to label if <code>b > a</code>	<code>jg exit</code>
<code>jge label</code>	Jump to label if <code>b ≥ a</code>	<code>jge format_disk</code>
<code>jl label</code>	Jump to label if <code>b < a</code>	<code>jl error</code>
<code>jle label</code>	Jump to label if <code>b ≤ a</code>	<code>jle finish</code>
<code>test reg,imm</code>	Bitwise compare of register and constant; should immediately precede the <code>jz</code> or <code>jnz</code> instructions	<code>test eax,0xffff</code>
<code>jz label</code>	Jump to label if bits were not set ("zero")	<code>jz looparound</code>
<code>jnz label</code>	Jump to label if bits were set ("not zero")	<code>jnz error</code>
<code>jmp label</code>	Unconditional relative jump	<code>jmp exit</code>
<code>jmp reg</code>	Unconditional absolute jump; arg is a register	<code>jmp eax</code>
Miscellaneous		
<code>nop</code>	No-op (opcode 0x90)	<code>nop</code>
<code>hlt</code>	Halt the CPU	<code>hlt</code>

Instructions with no memory references must include 'byte', 'word' or 'dword' size specifier.

Arguments to instructions: Note that it is not possible for **both** src and dest to be memory addresses.

Constant (decimal or hex): 10 or 0xff Fixed address: [200] or [0x1000+53]

Register: eax bl Dynamic address: [eax] or [esp+16]

32-bit registers: eax, ebx, ecx, edx, esi, edi, ebp, esp (points to first used location on top of stack)

16-bit registers: ax, bx, cx, dx, si, di, sp, bp

8-bit registers: al, ah, bl, bh, cl, ch, dl, dh